Eric Nalisnick

e.t.nalisnick@uva.nl   enalisnick.github.io

Academic Employment

Assistant Professor, Tenure Track   University of Amsterdam
Amsterdam Machine Learning Lab   September 2020 - present

Postdoctoral Research Associate   University of Cambridge
Supervisor: José Miguel Hernández-Lobato   September 2018 to September 2020

Graduate Student Researcher   University of California, Irvine
Supervisor: Padhraic Smyth   October 2013 to June 2018

Industrial Employment

Research Scientist   Google DeepMind
Part time, one day per week (20% FTE)   February 2019 to January 2020

Research Scientist Intern   Google DeepMind
Supervisor: Balaji Lakshminarayanan   Summer 2018

Applied Scientist Intern   Amazon
Supervisors: Vijai Mohan, Eiman Elnahrawy   Fall 2016

Research Intern   Twitter
Supervisor: Hugo Larochelle   Summer 2016

Research Intern   Microsoft
Supervisors: Rich Caruana, Nick Craswell   Summer 2015

Research Scientist Intern   Amazon
Supervisors: Vijai Mohan, Rahul Bhagat   Summer 2014

Education

Ph.D. Computer Science   2013 - 2018
University of California, Irvine

M.S. Computer Science   2012 - 2013
Lehigh University

B.S. Computer Science & English Literature   2008 - 2012
Lehigh University

Academic Honors and Awards

Veni Laureate, Dutch Research Council (NWO) 2021

ELLIS Scholar, *European Lab for Learning and Intelligent Systems Society* 2021

Top / Best Reviewer NeurIPS 2017, ICML 2019, ICML 2020

Research Funding

PRINCIPAL INVESTIGATOR

Continual Learning under Human Guidance June 2022 - 2026
€280,000
Veni, Talent Programme, Dutch Research Council (NWO): Science Domain
Single Principal Investigator, Acceptance Rate: 16%.

CO-INVESTIGATOR

UvA-Bosch Delta Lab November 2021 - 2026
Gift funding for 10 PhD students from the Bosch Group.
Role: Supervisor for 3 PhD students.
PIs: Theo Gevers, Jan-Willem van de Meent.

Hybrid Intelligence Centre January 2020 - 2030
Gravitation Programme, Dutch Research Council (NWO)
Role: Co-supervisor for 2 PhD students.
PIs: F. v. Harmelen, C. Jonker, M. d. Rijke, R. Verbrugge, P. Vossen, M. Welling.

Academic Supervision

PHD CANDIDATES

Mona Schirmer University of Amsterdam, 2022 -

Dharmesh Tailor University of Amsterdam, 2021 -

Saba Amiri (with Adam Belloum and Sander Klous) University of Amsterdam, 2021 -

Putra Manggala (with Holger Hoos) University of Amsterdam, 2021 -

Urja Khurana (with Antske Fokkens) Vrije University Amsterdam, 2020 -

Mrinank Sharma (with Tom Rainforth and Yee Whye Teh) University of Oxford, 2020 -

MASTERS STUDENTS

Nils Lehmann University of Amsterdam, 2021 -

Shuai Wang University of Amsterdam, 2021 -

Rajeev Verma University of Amsterdam, 2021 -

Arsen Sheverdin University of Amsterdam, 2021 -

Daniël Nobbe   University of Amsterdam, 2021

Teaching

Introduction to Machine Learning (*Leren*), University of Amsterdam   2020 - present
∼ 180 Undergraduate Students, ∼ 12 Teaching Assistants

Bayesian Deep Learning Module, Deep Learning II, University of Amsterdam   2022
123 Graduate Students, 5 Teaching Assistants

Professional Service

ORGANIZATION

Anomaly Detection for Scientific Discovery   2021 - present

Tractable Probabilistic Modeling, UAI Workshop   2022

Bayesian Deep Learning, NeurIPS Workshop   2021

Bayesian Deep Learning, ELLIS Workshop / NeurIPS Meetup   2020

Bayesian Deep Learning, NeurIPS Workshop   2019

AREA CHAIR / SENIOR PROGRAM COMMITTEE

Uncertainty in Artificial Intelligence (UAI)   2021 - present

International Conference on Machine Learning (ICML)   2022

Artificial Intelligence and Statistics (AIStats)   2022

International Conference on Learning Representations (ICLR)   2021

EDITORIAL BOARD

Editor, Probabilistic Methods for Deep Learning, Special Issue of *Entropy*   2021

JOURNAL REVIEWING

American Statistical Association (JASA)   2020 - present

Machine Learning Research (JMLR)   2018 - present

Advances in Statistical Analysis   2020 - present

CONFERENCE REVIEWING

Neural Information Processing Systems (NeurIPS)   2016 - present

International Conference on Learning Representations (ICLR)   2018 - present

International Conference on Machine Learning (ICML)   2018 - present

Artificial Intelligence and Statistics (AIStats)   2019 - present

Uncertainty in Artificial Intelligence (UAI)	2019 - present
Hybrid Human-Artificial Intelligence (HHAI)	2022
Association for the Advancement of Artificial Intelligence (AAAI)	2020 - 2021
International Joint Conference on Artificial Intelligence (IJCAI)	2019

Departmental / Institute Service

Education program committee for Bachelors and Masters of AI	2021 - present
Hiring committee for one AMLab / ELLIS unit faculty position	2021
Hiring committee for two AMLab faculty positions	2021

Publications

* Denotes equal contribution

JOURNAL ARTICLES

1. G. Papamakarios*, **E. Nalisnick***, D. J. Rezende, S. Mohamed, and B. Lakshminarayanan. Normalizing Flows for Probabilistic Modeling and Inference. *Journal of Machine Learning Research* (JMLR), 2021.

CONFERENCE ARTICLES

2. R. Verma and **E. Nalisnick**. Calibrated Learning to Defer with One-vs-All Classifiers. In *Proceedings of the 39th International Conference on Machine Learning* (ICML), 2022.

3. J. Antoran, J. U. Allingham, D. Janz, E. Daxberger, **E. Nalisnick**, and J. M. Hernández-Lobato. Adapting the Linearised Laplace Model Evidence for Modern Deep Learning. In *Proceedings of the 39th International Conference on Machine Learning* (ICML), 2022.

4. E. Daxberger, **E. Nalisnick***, J. U. Allingham*, J. Antoran*, and J. M. Hernández-Lobato. Expressive yet Tractable Bayesian Deep Learning via Subnetwork Inference. In *Proceedings of the 38th International Conference on Machine Learning* (ICML), 2021.

5. **E. Nalisnick**, J. Gordon, and J. M. Hernández-Lobato. Predictive Complexity Priors. In *Proceedings of the 24th International Conference on Artificial Intelligence and Statistics* (AISTATS), 2021.

6. R. Pinsler, J. Gordon, **E. Nalisnick**, and J. M. Hernández-Lobato. Bayesian Batch Active Learning as Sparse Subset Approximation. In *Advances in Neural Information Processing Systems* (NeurIPS), 2019.

7. **E. Nalisnick**, J. M. Hernández-Lobato, and P. Smyth. Dropout as a Structured Shrinkage Prior. In *Proceedings of the 36th International Conference on Machine Learning* (ICML), 2019.

8. **E. Nalisnick***, A. Matsukawa*, Y. W. Teh, D. Gorur, and B. Lakshminarayanan. Hybrid Models with Deep and Invertible Features. In *Proceedings of the 36th International Conference on Machine Learning* (ICML), 2019.

9. **E. Nalisnick**, A. Matsukawa, Y. W. Teh, D. Gorur, and B. Lakshminarayanan. Do Deep Generative Models Know What They Don't Know? In *Proceedings of the 7th International Conference on Learning Representations* (ICLR), 2019.

10. D. Ji, **E. Nalisnick**, Y. Qian, R. Scheuermann, and P. Smyth. Bayesian Trees for Automated Cytometry Data Analysis. In *Proceedings of Machine Learning for Healthcare* (MLHC), 2018.

11. **E. Nalisnick** and P. Smyth. Learning Priors for Invariance. In *Proceedings of the 21st International Conference on Artificial Intelligence and Statistics* (AISTATS), 2018.

12. **E. Nalisnick** and P. Smyth. Learning Approximately Objective Priors. In *Proceedings of the 33rd Conference on Uncertainty in Artificial Intelligence* (UAI), 2017.

13. **E. Nalisnick** and P. Smyth. Stick-Breaking Variational Autoencoders. In *Proceedings of the 5th International Conference on Learning Representations* (ICLR), 2017.

14. **E. Nalisnick**, B. Mitra, N. Craswell, and R. Caruana. Improving Document Ranking with Dual Word Embeddings. In *Proceedings of the 25th World Wide Web Conference* (WWW), 2016.

15. **E. Nalisnick** and H. Baird. Character-to-Character Sentiment Analysis in Shakespeare's Plays. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics* (ACL), 2013.

16. **E. Nalisnick** and H. Baird. Extracting Sentiment Networks from Shakespeare's Plays. In *Proceedings of the 12th International Conference on Document Analysis and Recognition* (ICDAR), 2013.

Peer-Reviewed Workshop Articles

17. U. Khurana, I. Vermeulen, **E. Nalisnick**, and A. Fokkens. Hate Speech Criteria: A Modular Approach to Task-Specific Hate Speech Definitions. *Workshop on Online Abuse and Harms*, NAACL 2022.

18. S. Amiri, A. Belloum, **E. Nalisnick**, S. Klous, and L. Gommans. On the Impact of Non-IID Data on the Performance and Fairness of Differentially Private Federated Learning. *Dependable and Secure Machine Learning*, DSN 2022.

19. J. Antoran, J. U. Allingham, D. Janz, E. Daxberger, **E. Nalisnick**, and J. M. Hernández-Lobato. Linearised Laplace Inference in Networks with Normalisation Layers and the Neural g-Prior. *Symposium on Advances in Approximate Bayesian Inference*, 2022.

20. U. Khurana, **E. Nalisnick**, and A. Fokkens. How Emotionally Stable is ALBERT? Testing Robustness with Stochastic Weight Averaging on a Sentiment Analysis Task. *Evaluation and Comparison of NLP Systems*, EMNLP 2021.

21. P. Manggala, H. Hoos, and **E. Nalisnick**. Bayesian Regression from Multiple Sources of Weak Supervision. *Machine Learning for Data*, ICML 2021.

22. Y. Zhang and **E. Nalisnick**. On the Inconsistency of Bayesian Inference for Misspecified Neural Networks. *Symposium on Advances in Approximate Bayesian Inference*, 2021.

23. E. Daxberger, **E. Nalisnick***, J. U. Allingham*, J. Antoran*, and J. M. Hernández-Lobato. Expressive yet Tractable Bayesian Deep Learning via Subnetwork Inference. *Symposium on Advances in Approximate Bayesian Inference*, 2021.

24. **E. Nalisnick**, J. Gordon, and J. M. Hernández-Lobato. Predictive Complexity Priors. *Uncertainty & Robustness in Deep Learning*, ICML 2020.

25. **E. Nalisnick**, A. Matsukawa, Y. W. Teh, and B. Lakshminarayanan. Detecting Out-of-Distribution Inputs to Deep Generative Models Using Typicality. *Bayesian Deep Learning*, NeurIPS 2019.

26. **E. Nalisnick** and J. M. Hernández-Lobato. Automatic Depth Determination for Bayesian ResNets. *Bayesian Deep Learning*, NeurIPS 2018.

27. **E. Nalisnick**, A. Matsukawa, Y.W. Teh, D. Gorur, and B. Lakshminarayanan. Do Deep Generative Models Know What They Don't Know? *Bayesian Deep Learning*, NeurIPS 2018.

28. **E. Nalisnick***, A. Matsukawa*, Y.W. Teh, D. Gorur, and B. Lakshminarayanan. Hybrid Models with Deep and Invertible Features. *Bayesian Deep Learning*, NeurIPS 2018.

29. O. Rybakov, V. Mohan, A. Misra, S. LeGrand, R. Joseph, K. Chung, S. Singh, Q. You, **E. Nalisnick**, L. Dirac, and R. Luo. The Effectiveness of a Two-Layer Neural Network for Recommendations. Workshop Track, ICLR 2018.

30. D. Ji, **E. Nalisnick**, and P. Smyth. Mondrian Processes for Flow Cytometry Analysis. *Machine Learning for Health*, NeurIPS 2017.

31. **E. Nalisnick** and P. Smyth. Variational Inference with Stein Mixtures. *Advances in Approximate Bayesian Inference*, NIPS 2017.

32. **E. Nalisnick** and P. Smyth. The Amortized Bootstrap. *Implicit Models*, ICML 2017.

33. **E. Nalisnick** and P. Smyth. Variational Reference Priors. Workshop Track, ICLR 2017.

34. **E. Nalisnick**, L. Hertel, and P. Smyth. Approximate Inference for Deep Latent Gaussian Mixtures. *Bayesian Deep Learning*, NeurIPS 2016.

35. **E. Nalisnick** and P. Smyth. Nonparametric Deep Generative Models with Stick-Breaking Priors. *Data-Efficient Machine Learning*, ICML 2016.

36. J. Park, M. Blume-Kohout, R. Krestel, **E. Nalisnick**, and P. Smyth. Analyzing NIH Funding Patterns over Time with Statistical Text Analysis. *Scholarly Big Data*, AAAI 2016.

Theses

1. **E. Nalisnick**. On Priors for Bayesian Neural Networks. *Doctoral Dissertation*, University of California, Irvine, 2018.

2. **E. Nalisnick**. Automatic Methods for Tracking Sentiment Dynamics in Plays. *Master's Thesis*, Lehigh University, 2013.

3. **E. Nalisnick**. A Combinatorial Explanation for a Conjecture of Fomin and Zelevinsky. *Honors Thesis*, Lehigh University, 2012.

Patents

1. E. M. H. Elnahrawy, V. Mohan, and **E. Nalisnick**. Generation and Use of Model Parameters in Cold-Start Scenarios. U.S. Patent Number 10,726,334. 28 July 2020.

Invited Talks

1. Towards Informative Priors for Bayesian Deep Learning, DAGSTUHL SEMINAR — 2022
2. On the Calibration of Learning-to-Defer Systems, UNIV. OF MANCHESTER, STAT. SEMINAR — 2022
3. Predictive Complexity Priors, UNIVERSITY OF EDINBURGH, ANC SEMINAR — 2021
4. Predictive Complexity Priors, GENERATIVE MODELS AND UNCERTAINTY QUANTIFICATION — 2021
5. Predictive Complexity Priors, IMPERIAL COLLEGE STATISTICS SEMINAR — 2021
6. Predictive Complexity Priors, ALAN TURING INSTITUTE — 2020
7. Detecting Distribution Shift with Deep Generative Models, SYDNEY ML MEETUP — 2020
8. Detecting Distribution Shift with Deep Generative Models, INNF+, ICML WORKSHOP — 2020
9. Building and Critiquing Models for Probabilistic Deep Learning, GATSBY UNIT, UCL — 2020
10. Building and Critiquing Models for Probabilistic Deep Learning, CARNEGIE MELLON UNIV. — 2020
11. Building and Critiquing Models for Probabilistic Deep Learning, UNIV. OF NORTH CAROLINA — 2020
12. Deep Learning & Statistics: Bridging the Gap with Prob. Structure, UNIV. OF AMSTERDAM — 2020
13. Deep Learning & Statistics: Bridging the Gap with Prob. Structure, UC SANTA BARBARA — 2020
14. Deep Learning Under Covariate Shift, UCI AI/ML SEMINAR — 2019
15. Normalizing Flows for Tractable Probabilistic Modeling and Inference, T-PRIME, NEURIPS — 2019
16. Deep Learning: A Synthesis from Probabilistic Foundations, RAND CORP STATS. SEMINAR — 2019
17. Evaluating Deep Generative Models on Out-of-Distribution Inputs, OXFORD STATS. SEMINAR — 2019
18. Do Deep Generative Models Know What They Don't Know?, CAMAIML (MSR CAMBRIDGE) — 2019
19. Do Deep Generative Models Know What They Don't Know?, CAMBRIDGE LTL SEMINAR — 2019
20. Structured Shrinkage Priors for Neural Networks, IMPERIAL COLLEGE STATISTICS SEMINAR — 2018
21. Deep Learning: A Synthesis from Probabilistic Foundations, UCI STATISTICS SEMINAR — 2018
22. Approximate Inference for Frequentist Uncertainty Estimation, SOCAL ML SYMPOSIUM — 2017
23. Deep Generative Models with Stick-Breaking Priors, UCI AI/ML SEMINAR — 2017
24. Alternative Priors for Deep Generative Models, OPENAI — 2017